

# Reinforcement Learning: From Q Learning to Deep Q-Learning

Prasun De (Roll: MB2117)

ISI Kolkata

July 25, 2023

# Reinforcement Learning

- Learning to map situations to actions in order to maximize a numerical reward
- Learner must **discover** which actions to take by trying them out ( trial-and-error )
- Learner can face **delayed reward**: consequences of a delayed action are not immediately observable or received
- Learner faces **sequential data and interactions**, where the learner's actions influence the subsequent states and rewards it receives

# Markov Decision Process (MDP)




Figure: Agent–Environment Interaction

Assumptions:

- Agent gets to observe the state
- **Markov Property:**

$$\begin{aligned} P(s_{t+1} = s', r_{t+1} = r | s_t, a_t, r_t, s_{t-1}, a_{t-1}, \dots, r_1, s_0, a_0) \\ = P(s_{t+1} = s', r_{t+1} = r | s_t, a_t) \end{aligned}$$

- **Self-Driving Cars:** Q-Learning for Lane changing [2]
- **Energy Conservation:** Deepmind uses RL agents to achieve 40% decrease in energy consumption of data centers
- **Financial Applications:** Hedging portfolio in presence of transaction costs etc [1]
- **Healthcare:** Discovery and generation of optimal DTRs for chronic diseases [8]
- **Recommendation Engines:** Personalized User Recommendations [3]

# MDP Setup

MDP is defined by:

- Set of States  $S$
- Set of Actions  $A$
- Transition function  $P(s'|s, a)$
- Reward function  $R(s, a, s')$
- Start state  $s_0$
- Discount factor  $\gamma$
- Horizon  $H$

**Goal:** Given an  $\text{MDP}(S, A, P, R, \gamma, H)$  obtain **optimal policy**  $\pi^*$  for

$$\max_{\pi} E \left[ \sum_{t=0}^H \gamma^t R(S_t, A_t, S_{t+1}) | \pi \right]$$

# Exact Method: Value Iteration

- **Optimal Value Function**  $V_H^*$  for horizon  $H$

$$V_H^*(s) := \max_{\pi} E \left[ \sum_{t=0}^H \gamma^t R(S_t, A_t, S_{t+1}) \mid \pi, s_0 = s \right]$$

= sum of discounted rewards starting from state  $s$   
and acting optimally

- We can say that  $V_0^*(s) = 0 \forall s$ .

$$\begin{aligned} V_1^*(s) &= \text{optimal value for state } s \text{ when } H = 0 \\ &= \max_a \sum_{s'} P(s'|s, a) \cdot (R(s, a, s') + \gamma V_0^*(s')) \end{aligned}$$

- In general, we can say:

$$\begin{aligned} V_k^*(s) &= \text{optimal value for state } s \text{ when } H = k \\ &= \max_a \sum_{s'} P(s'|s, a) \cdot (R(s, a, s') + \gamma V_{k-1}^*(s')) \end{aligned}$$

# Value Iteration Convergence-I

Algorithm:

Start with  $V_0^*(s) = 0$  for all  $s$ .

For  $k = 1, \dots, H$ :

For all states  $s$  in  $S$ :

$$V_k^*(s) \leftarrow \max_a \sum_{s'} P(s'|s, a) (R(s, a, s') + \gamma V_{k-1}^*(s'))$$

$$\pi_k^*(s) \leftarrow \arg \max_a \sum_{s'} P(s'|s, a) (R(s, a, s') + \gamma V_{k-1}^*(s'))$$

This is called a **value update** or **Bellman update/back-up**

Figure: Value Iteration Algorithm

**Result:** Value iteration converges. At convergence, we obtain the optimal value  $V^*$  for infinite horizon problem and it satisfies:

$$V^*(s) = \max_a \sum_{s'} P(s'|s, a) \cdot (R(s, a, s') + \gamma V^*(s'))$$

# Value Iteration convergence-II

- Infinite horizon policy is stationary: optimal action at a state  $s$  is the same action at all times
- Recall definitions of  $V^*(s)$  and  $V_H^*(s)$  for a state  $s$ .
- Additional reward collected over  $t = H + 1, H + 2, \dots$

$$\begin{aligned}\gamma^{H+1}R(s_{H+1}) + \gamma^{H+2}R(s_{H+2}) + \dots &\leq \gamma^{H+1}R_{\max} + \gamma^{H+2}R_{\max} + \dots \\ &= \frac{\gamma^{H+1}}{1 - \gamma}R_{\max} \rightarrow 0\end{aligned}$$

Intuition for  $V_H^* \rightarrow V^*$  as  $H \rightarrow \infty$

- Proof involves contractions of the max-norm



- $Q^*(s, a)$  is the expected utility started in  $s$ , taking action  $a$  and (thereafter) acting optimally
- **Q-value iteration** has a somewhat similar form to the Value function iteration

$$Q_{k+1}^*(s, a) = \sum_{s'} P(s'|s, a)(R(s, a, s') + \gamma \max_{a'} Q_k^*(s', a'))$$

- **Bellman Equation:**

$$Q^*(s, a) = \sum_{s'} P(s'|s, a)(R(s, a, s') + \gamma \max_{a'} Q^*(s', a'))$$

- No need to keep track of policy and value now !

# (Tabular) Q-Learning

## ■ Q-value iteration

$$Q_{k+1}^*(s, a) = \sum_{s'} P(s'|s, a) (R(s, a, s') + \gamma \max_{a'} Q_k^*(s', a'))$$

## ■ Rewriting it as an expectation:

$$Q_{k+1}^*(s, a) = E_{s' \sim P(s'|s, a)} \left[ R(s, a, s') + \gamma \max_{a'} Q_k^*(s', a') \right]$$

## ■ Tabular Q-Learning: Replace expectation by samples.

- For state-action pair  $(s, a)$ , receive  $s' \sim P(s'|s, a)$
- Consider **old estimate**  $Q_k(s, a)$
- Obtain **new sample estimate**:

$$\text{target}(s') = R(s, a, s') + \gamma \max_{a'} Q_k(s', a')$$

- Incorporate new estimate into running average

$$Q_{k+1}(s, a) = (1 - \alpha) Q_k(s, a) + \alpha \text{target}(s')$$

# Tabular Q-Learning Algorithm

Algorithm:

```
Start with  $Q_0(s, a)$  for all  $s, a$ .  
Get initial state  $s$   
For  $k = 1, 2, \dots$  till convergence  
    Sample action  $a$ , get next state  $s'$   
    If  $s'$  is terminal:  
        target =  $R(s, a, s')$   
        Sample new initial state  $s'$   
    else:  
        target =  $R(s, a, s') + \gamma \max_{a'} Q_k(s', a')$   
         $Q_{k+1}(s, a) \leftarrow (1 - \alpha)Q_k(s, a) + \alpha [\text{target}]$   
         $s \leftarrow s'$ 
```

Figure: Tabular Q-Learning Algorithm

We can sample actions as follows:


- Choose  $a$  randomly
- Choose  $a$  that maximizes  $Q_k(s, a)$  greedily
- $\epsilon$ -Greedy: Choosing random action w.p.  $\epsilon$  and greedily w.p.  $1 - \epsilon$

# Q-Learning Convergence


- **Q-Learning Converges** to the optimal policy (under some weak conditions)
- **Off-policy learning**
- Explore enough  $(s, a)$  pairs and also exploit the learnt  $Q$  – values
- Eventually decrease learning rate gradually
- **Convergence guarantee** for bounded reward and finite state-action space [7]:
  - $n^i(s, a)$  is the index of  $i$ -th time that action  $a$  is tried in state  $s$
  - $|r_n| \leq R_{\max}$
  - $\alpha_n \in [0, 1]$
  - $\sum_{i=1}^{\infty} \alpha_{n^i(s, a)} = \infty$
  - $\sum_{i=1}^{\infty} \alpha_{n^i(s, a)}^2 < \infty$

Then,  $Q_n(s, a) \rightarrow Q^*(s, a) \forall s, a$

# Scaling Tabular Methods



Gridworld  
 $10^1$



Tetris  
 $10^{60}$



Atari  
 $10^{308}$  (ram)  $10^{16992}$  (pixels)

Figure: Size of State Space

# Challenges with scaling

- Basic Q-Learning keeps a table of Q-values.
- In realistic situation, cannot possibly learn about every single state.
  - Too many states to visit all during training
  - Too many state-action pairs to hold Q-table in memory
- We would instead want to generalize:
  - Learn about small number of training states from experience
  - Generalize the experience to new similar situations

# Approximate Q-Learning

- We have a parameterized Q-function  $Q_\theta(s, a)$  instead of a table.

- $Q_\theta(s, a) = \sum_{k=0}^n \theta_k f_k(s, a)$  or

- Can be a neural network

- **Learning Rule:** We can try to change the

- $\text{target}(s') = R(s, a, s') + \gamma \max_{a'} Q_{\theta_k}(s', a')$

- Gradient update on  $\theta$ :

- $$\theta_{k+1} = \theta_k - \alpha \nabla_\theta \left[ \frac{1}{2} (Q_\theta(s, a) - \text{target}(s'))^2 \right] \Big|_{\theta=\theta_k}$$

# Connection to Tabular Q-Learning

- Suppose  $\theta \in \mathbb{R}^{|S| \times |A|}$  and  $Q_\theta(s, a) = \theta_{sa}$

$$\begin{aligned} & \nabla_{\theta_{sa}} \left[ \frac{1}{2} (Q_{\theta_{sa}}(s, a) - \text{target}(s'))^2 \right] \\ &= \nabla_{\theta_{sa}} \left[ \frac{1}{2} (\theta_{sa} - \text{target}(s'))^2 \right] \\ &= \theta_{sa} - \text{target}(s') \end{aligned}$$

- Now, we may plug it into our update:

$$\begin{aligned} \theta_{sa} &\leftarrow \theta_{sa} - \alpha(\theta_{sa} - \text{target}(s')) \\ &= (1 - \alpha)\theta_{sa} + \alpha[\text{target}(s')] \end{aligned}$$

- Compare with tabular Q-Learning Update:

$$Q_{k+1}(s, a) \leftarrow (1 - \alpha)Q_k(s, a) + \alpha[\text{target}(s')]$$

- Deep RL essentially uses powerful function approximators (eg; CNN) to represent Q-values ( **with some care !** )



# Approximate Q-Learning

Algorithm:

Start with  $Q_0(s, a)$  for all  $s, a$ .

Get initial state  $s$

For  $k = 1, 2, \dots$  till convergence

Sample action  $a$ , get next state  $s'$

If  $s'$  is terminal:

$$\text{target} = R(s, a, s')$$

Sample new initial state  $s'$

else:

$$\text{target} = R(s, a, s') + \gamma \max_{a'} Q_k(s', a')$$

$$\theta_{k+1} \leftarrow \theta_k - \alpha \nabla_{\theta} \mathbb{E}_{s' \sim P(s'|s,a)} [(Q_{\theta}(s, a) - \text{target}(s'))^2] \big|_{\theta=\theta_k}$$

$$s \leftarrow s'$$

Chasing a nonstationary target!

Updates are correlated within a trajectory!

Figure: Approx. Q-Learning Algorithm

# Deep Q-Networks

- The high-level idea is to make Q-Learning look like supervised learning
- Two essential ideas to stabilize training:
  - **Experience Replay Buffer** [4]
  - Previously used for better data efficiency
  - Makes data distribution more stationary
  - Use an **older set of weights to compute targets**
  - Keeps target function from changing too quickly

$$L_i(\theta_i) = \mathbb{E}_{s,a,s',r \sim D} \left( \underbrace{r + \gamma \max_{a'} Q(s', a'; \theta_i^-)}_{\text{target}} - Q(s, a; \theta_i) \right)^2$$

Figure: DQN update

# Replay Buffer

- Most recent  $k$  transitions  $e_t = (s_t, a_t, r_t, s_{t+1})$  are stored in a **replay buffer**  $D_T = \{e_1, e_2, \dots, e_T\}$
- Sample uniformly a batch of  $N$  transitions to update the Q-network
- Helps in improving data efficiency, reducing sample correlations




Figure: Experience Replay

# Target Network Intuition

- Changing the value of one action will change the value of other actions and similar states.
- The network can end up chasing its own tail because of bootstrapping.

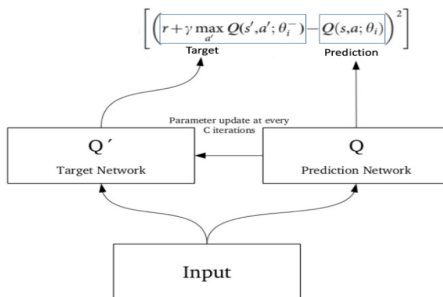


Figure: Target network

# DQN Algorithm

**Algorithm 1: deep Q-learning with experience replay.**

Initialize replay memory  $D$  to capacity  $N$

Initialize action-value function  $Q$  with random weights  $\theta$

Initialize target action-value function  $\hat{Q}$  with weights  $\theta^- = \theta$

**For** episode = 1,  $M$  **do**

    Initialize sequence  $s_1 = \{x_1\}$  and preprocessed sequence  $\phi_1 = \phi(s_1)$

**For**  $t = 1, T$  **do**

        With probability  $\varepsilon$  select a random action  $a_t$

        otherwise select  $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

        Execute action  $a_t$  in emulator and observe reward  $r_t$  and image  $x_{t+1}$

        Set  $s_{t+1} = s_t, a_t, x_{t+1}$  and preprocess  $\phi_{t+1} = \phi(s_{t+1})$

        Store transition  $(\phi_t, a_t, r_t, \phi_{t+1})$  in  $D$

        Sample random minibatch of transitions  $(\phi_j, a_j, r_j, \phi_{j+1})$  from  $D$

        Set  $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

        Perform a gradient descent step on  $(y_j - Q(\phi_j, a_j; \theta))^2$  with respect to the network parameters  $\theta$

        Every  $C$  steps reset  $\hat{Q} = Q$

**End For**

**End For**

Figure: Target network

# DQN on Atari




Figure: Some Atari Games


- 49 ATARI 2600 games.
- From pixels to actions.
- The change in score is the reward.
- Same algorithm.
- Same function approximator, w/ 3M free parameters.
- Same hyperparameters.
- Roughly human-level performance on 29 out of 49 games.

# Double DQN

- There is an **upward bias** in  $\max_a Q(s, a; \theta)$  [6]
- DQN maintains two sets of weight  $\theta$  and  $\theta^-$ , so reduce bias by using:
  - $\theta$  to **select** best action
  - $\theta^-$  to **evaluate** best action
- **Double DQN** [6] loss:

$$L_i(\theta_i) = \mathbb{E}_{s,a,s',r \sim D} \left( r + \gamma Q(s', \arg \max_{a'} Q(s', a'; \theta_i^-); \theta_i^-) - Q(s, a; \theta_i) \right)^2$$

Figure: Double DQN Loss



# Prioritized Experience Replay

- Replaying all transitions with equal probability is highly suboptimal.
- Replay transitions in proportion to absolute Bellman error [5]:

$$|r + \gamma \max_{a'} Q(s', a'; \theta^-) - Q(s, a, \theta)|$$

- Generally leads to faster learning




Figure: PER



Thank You !

- [1] Hans Bühler et al. *Deep Hedging*. 2018. arXiv: 1802.03042 [q-fin.CP].
- [2] B Ravi Kiran et al. “Deep reinforcement learning for autonomous driving: A survey”. In: *IEEE Transactions on Intelligent Transportation Systems* 23.6 (2021), pp. 4909–4926.
- [3] Lihong Li et al. “A contextual-bandit approach to personalized news article recommendation”. In: *Proceedings of the 19th international conference on World wide web*. ACM, Apr. 2010. DOI: 10.1145/1772690.1772758. URL: <https://doi.org/10.1145/1772690.1772758>.
- [4] Long-Ji Lin. *Reinforcement learning for robots using neural networks*. Carnegie Mellon University, 1992.
- [5] Tom Schaul et al. *Prioritized Experience Replay*. 2016. arXiv: 1511.05952 [cs.LG].
- [6] Hado Van Hasselt, Arthur Guez, and David Silver. “Deep reinforcement learning with double q-learning”. In:

- [7] Christopher JCH Watkins and Peter Dayan. “Q-learning”. In: *Machine learning* 8 (1992), pp. 279–292.
- [8] Chao Yu, Jiming Liu, and Shamim Nemati. *Reinforcement Learning in Healthcare: A Survey*. 2020. [arXiv: 1908.08796 \[cs.LG\]](#).